

# 聊聊PostgreSQL表膨胀

作者	时间	QQ技术交流群
<a href="mailto:perrynzhou@gmail.com">perrynzhou@gmail.com</a>	2022/08/20	672152841



存储内核技术交流

微信扫描二维码，关注我的公众号



开源存储问题解答社区:<https://github.com/perrynzhou/deep-dive-storage-in-china>

## PostgreSQL Basic

- PG中的MVCC(多版本并发)设计目的是读不阻塞写。PG中的所有的insert和update操作都是创建新的一行数据；update和delete都不是立即删除旧版本无用的数据。tuple是否可见是由snapshot决定。
- PG中追踪每个表的Block可见性是通过表的vm文件。Table或者Index的可用空间管理是通过表或者索引的fsm文件管理，它是一个2级的binary tree，最底层存储了每个page可用空间，最上层聚合最低层的信息。

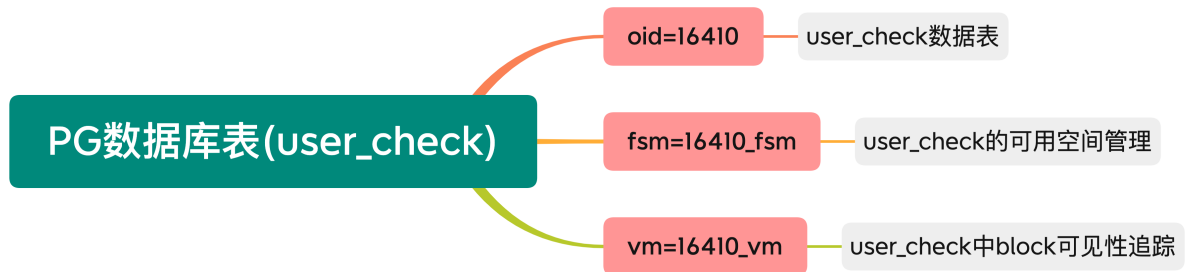
```
perryn_demo=> select oid,datname from pg_database where datname='perryn_demo';
 oid | datname 
-----+-----
 16394 | perryn_demo
(1 row)

perryn_demo=> select oid,relname from pg_class where relname='user_check';
 oid | relname 
-----+-----
 16410 | user_check
(1 row)
```

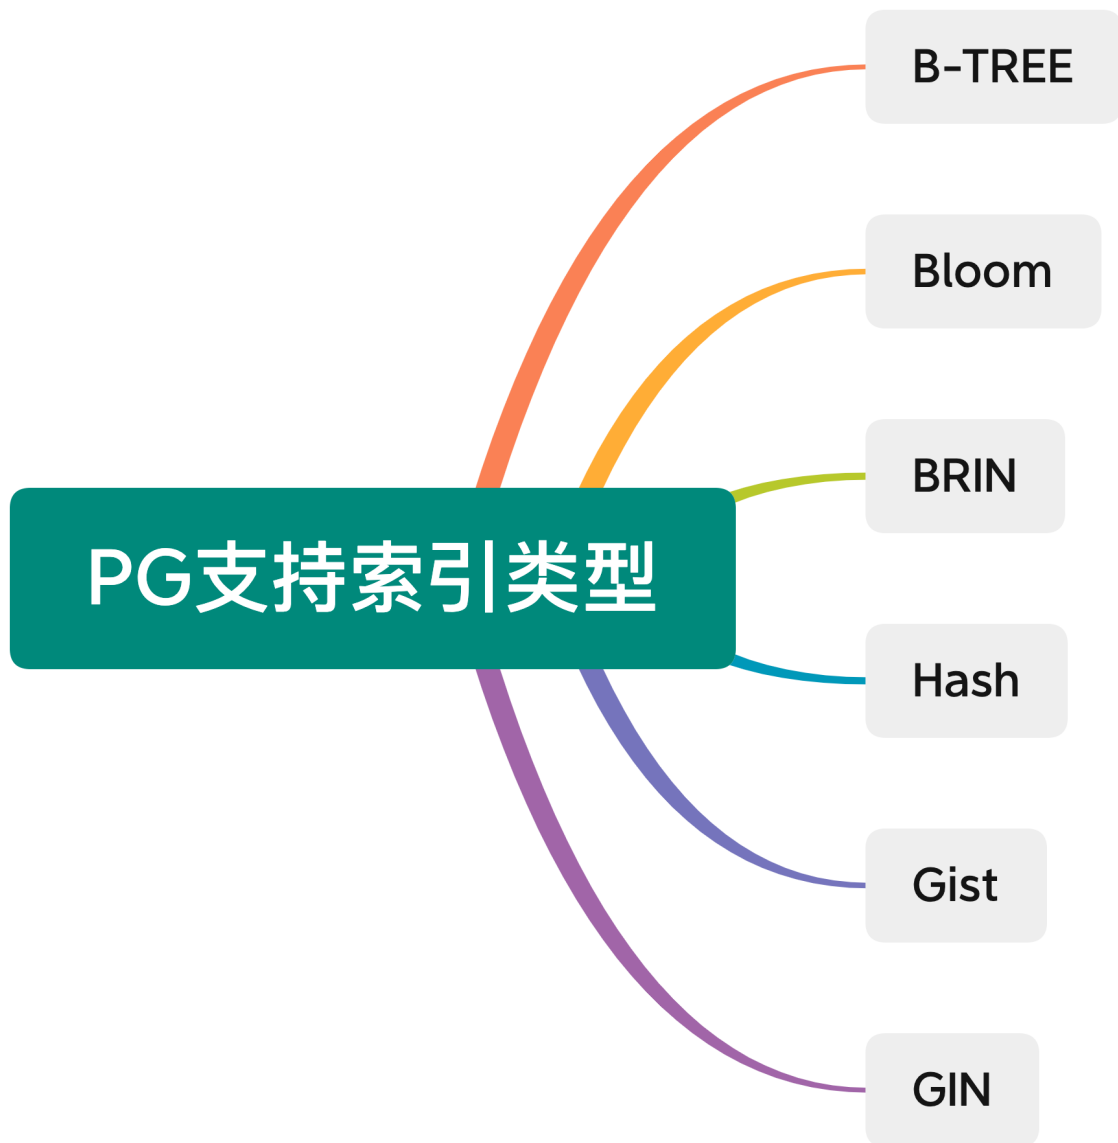
```

[perrynzhou@local-dev ~/Debug/pg_home/base]$ cd 16394/
[perrynzhou@local-dev ~/Debug/pg_home/base/16394]$ find ./ |grep 16410
./16410
./16410_vm
./16410_fsm
[perrynzhou@local-dev ~/Debug/pg_home/base/16394]$ ls -l -lh |grep 16410
-rw----- 1 perrynzhou perrynzhou 41M Aug 30 11:07 16410
-rw----- 1 perrynzhou perrynzhou 32K Aug 30 11:05 16410_fsm
-rw----- 1 perrynzhou perrynzhou 8.0K Aug 30 11:05 16410_vm

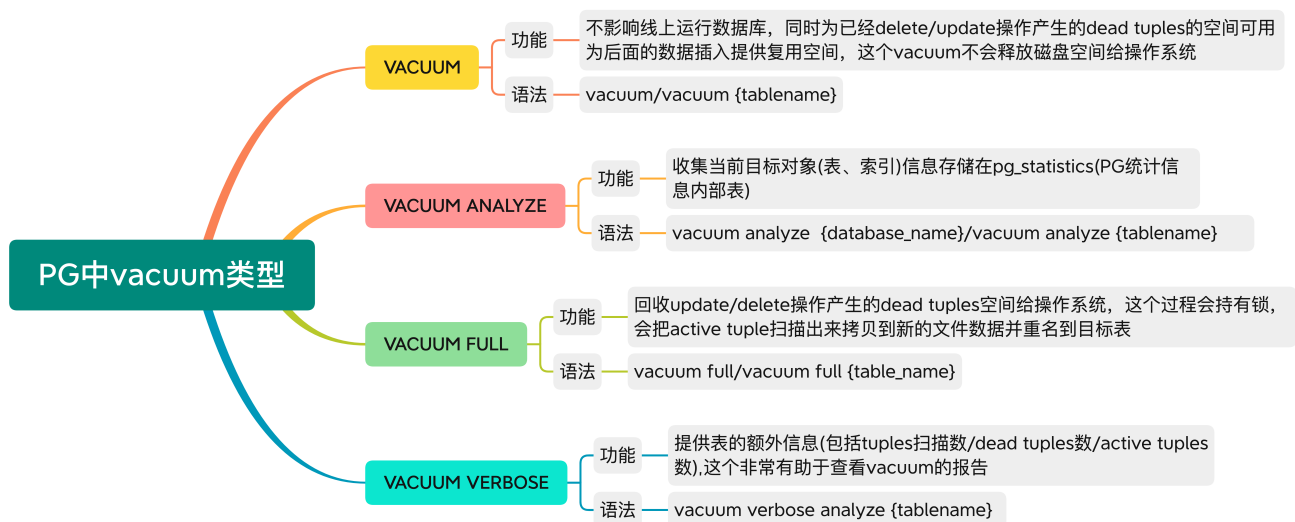
```



- PG目前支持多种索引类型，包括B-Tree、Hash、Gin、Gist、Brin、Bloom。



## PostgreSQL膨胀



- 膨胀在PG中表示表或者索引的大小大于实际数据的大小，其次表中每个block或者page的空间利用率低。当一个事务T1读取表的block B中A行数据时候，第二个事务T2去更新这

个表中Block B中A行数据；为了确保read事务不阻塞write事务，T2的write事务把更新后的A这一行数据写到新的空闲空间，而A这行数据依然在Block B中，这个就是dead tuple.所以在PG中，如果有非常多的update和delete,会产生非常多的dead tuples,这些dead tuples的集合就是PG中的膨胀。

- 针对PG中的膨胀问题是通过vacuum来解决，PG中的auto vacuum会阻塞read/write操作，手动的vacuum则不会阻塞。vacuum有三种类型，分别是普通的vacuum、vacuum analyze、vacuum full.

## 验证PostgreSQL膨胀

- OS版本

```
[perrynzhou@local-dev ~/Debug/pg_home]$ uname -a
Linux local-dev 4.18.0-348.7.1.el8_5.x86_64 #1 SMP Wed Dec 22 13:25:12 UTC
2021 x86_64 x86_64 x86_64 GNU/Linux
```

- PostgreSQL版本

```
[perrynzhou@local-dev ~/Debug/pg_home]$ psql --version
psql (PostgreSQL) 14.3
```

- 测试数据库和表信息

```
/*****设置测试数据库和登录用户*****/
[perrynzhou@local-dev ~/Debug]$ psql -d postgres
psql (14.3)
Type "help" for help.

// 创建测试数据库 perryn_demo
postgres=# create database perryn_demo;
CREATE DATABASE
// 创建perryn_demo数据库用户名称为perryn_demo
postgres=# CREATE USER perryn_demo WITH ENCRYPTED PASSWORD '123456';
CREATE ROLE

// 设置用户允许登录
postgres=# ALTER USER perryn_demo WITH login;
ALTER ROLE

// 授予perryn_demo数据库操作所有权限给用户perryn_demo
postgres=# grant all privileges on database perryn_demo to perryn_demo;
GRANT

/*****创建测试表和数据*****/
```

```
[perrynzhou@local-dev ~/Debug]$ psql -d perryn_demo -U perryn_demo
psql (14.3)
Type "help" for help.
perryn_demo=> create table user_check as select generate_series (1,10000)
as id, substr(md5(random()::text), 0, 255) as uuid, to_char(random() *
1000000, '099999') as code, substring(random()::varchar,3,8) as md5;
```

## 表的隐藏列

- PG中的隐藏列设计是为了MVCC功能设计，一个事务中的查询如何找到这个事务开启时候应该读取数据的版本。PG包含了这些隐藏列。是不同事务之间的数据版本判断的基础。是判断同一个事务内的其他命令导致的行版本变更是否可见

```
// 查询user_check表这个所有列(包括隐藏列)

perryn_demo=> drop table user_check;
DROP TABLE
perryn_demo=> create table user_check as select generate_series (1,5) as
id, substr(md5(random()::text), 0, 255) as uuid, to_char(random() *
1000000, '099999') as code, substring(random()::varchar,3,8) as md5;
SELECT 5
perryn_demo=> SELECT attrelid::regclass::text, attname, format_type
(atttypid, atttypmod) FROM pg_attribute WHERE
attrelid::regclass::text='user_check' ORDER BY attnum;
   attrelid   | attname   | format_type
-----+-----+-----
// tableoid是表的在PG内部唯一标识
user_check | tableoid | oid
// 删除事务中的命令标识
user_check | cmax     | cid
// 如果xmax 为0 ，表示数据没有被删除；如果不为0，则是删除这个数据的事务ID
user_check | xmax     | xid
// 插入事务中的命令标识
user_check | cmin     | cid
// xmin 是每个事务中数据插入时候的事务ID
user_check | xmin     | xid
user_check | ctid     | tid
user_check | id       | integer
user_check | uuid     | text
user_check | code     | text
user_check | md5      | text
(10 rows)
```

- xmin隐藏列表示数据插入时候的事务ID,xmax隐藏列表示数据删除/更改时候的事务ID.这次模拟是在会话A中初始化插入数据->会话B中更新数据->在回到会话A中查询数据来观察数据表是如何膨胀的。

```
// 禁用数据表的vacuum
ALTER TABLE ucheck SET (
    autovacuum_enabled = false, toast.autovacuum_enabled = false
);
```

```
// 会话A:查询当前的事务ID,事务ID=811
```

```
perryn_demo=> begin;
```

BEGIN

```
perryn_demo=> select txid_current();
```

txid\_current

-----

811

(1 row)

```
perryn_demo=> create table ucheck as select generate_series (1,3) as id,
substr(md5(random()::text), 0, 255) as uuid, to_char(random() * 1000000,
'099999') as code, substring(random()::varchar,3,8) as md5;
```

SELECT 3

```
perryn_demo=> select xmin,xmax,cmin,cmax,* from ucheck;
```

xmin	xmax	cmin	cmax	id	uuid	code
md5						

-----+

811	0	5	5	1	f12b88a762ec72f1885145b53148c79a
692255	60326622				

811	0	5	5	2	2c7cd94aaa74ce04ed7325a93acdeb03
290345	59971147				

811 | 0 | 5 | 5 | 3 | d3e703cd56522833b5fbadd1459b9aa0 |  
548640 | 96239513

(3 rows)

```
perryn_demo=*> commit;
```

COMMIT

```
// 会话B:更新uclick中字段, 事务ID=813
```

```
perryn_demo=> select txid_current();
```

txid\_current

-----

813

(1 row)

// 这里会话B中更新时候插入了2条数据, 会话A中原来旧版本数据依然存在

```
perryn_demo=> update ucheck set md5=substring(random()::varchar,3,8) where
id>=2;
```

UPDATE 2

```
perryn_demo=> select xmin,xmax,cmin,cmax,* from ucheck
```

```
;
```

xmin	xmax	cmin	cmax	id	uuid	code
md5						
811	0	5	5	1	f12b88a762ec72f1885145b53148c79a	
692255	60326622					
813	0	0	0	2	2c7cd94aaa74ce04ed7325a93acdeb03	
290345	66262299					
813	0	0	0	3	d3e703cd56522833b5fbadd1459b9aa0	
548640	29500328					

(3 rows)

// 会话A:再次查看ucheck表的数据, xmax事务ID是为更新的事务ID, 这里就造成了表的膨胀

```
perryn_demo=> begin;
```

```
BEGIN
```

```
perryn_demo=> select xmin,xmax,cmin,cmax,* from ucheck;
```

xmin	xmax	cmin	cmax	id	uuid	code
md5						
811	0	5	5	1	f12b88a762ec72f1885145b53148c79a	
692255	60326622					
811	0	5	5	2	2c7cd94aaa74ce04ed7325a93acdeb03	
290345	59971147					
811	0	5	5	3	d3e703cd56522833b5fbadd1459b9aa0	
548640	96239513					

(3 rows)

// 这里观察到xmax = 会话B中的事务ID

```
perryn_demo=> select xmin,xmax,cmin,cmax,* from ucheck;
```

xmin	xmax	cmin	cmax	id	uuid	code
md5						
811	0	5	5	1	f12b88a762ec72f1885145b53148c79a	
692255	60326622					
811	813	0	0	2	2c7cd94aaa74ce04ed7325a93acdeb03	
290345	59971147					
811	813	0	0	3	d3e703cd56522833b5fbadd1459b9aa0	
548640	96239513					

(3 rows)

```
perryn_demo=> commit;
```

```
COMMIT
```

## 分析膨胀表的空间

## pageinspect查看表的dead tuples

- 授权perryn\_demo为SUPERUSER

```
[perrynzhou@local-dev ~]$ psql -d postgres
psql (14.3)
Type "help" for help.

postgres=#
postgres=# ALTER ROLE perryn_demo SUPERUSER;
ALTER ROLE
```

- 查看表的dead tuples

```
[perrynzhou@local-dev ~]$ psql -U perryn_demo -d perryn_demo
psql (14.3)
Type "help" for help.

perryn_demo=# CREATE EXTENSION pageinspect;
CREATE EXTENSION

// t_xmax中的813都是dead tuples, 目前这个表已经被禁用auto vacuum
perryn_demo=# SELECT t_xmin, t_xmax, tuple_data_split('ucheck'::regclass,
t_data, t_infomask, t_infomask2, t_bits) FROM
heap_page_items(get_raw_page('ucheck', 0));
 t_xmin | t_xmax |
tuple_data_split
-----+-----+-----
-----+-----+-----
---
      811 |      0 |
{"\\x01000000", "\\x43663132623838613736326563373266313838353134356235333134
3863373961", "\\x1120363932323535", "\\x133630333236363232"}
      811 |     813 |
{"\\x02000000", "\\x433263376364393461613734636530346564373332356139336163
6465623033", "\\x1120323930333435", "\\x133539393731313437"}
      811 |     813 |
{"\\x03000000", "\\x43643365373033636435363532323833336235666261646431343539
6239616130", "\\x1120353438363430", "\\x133936323339353133"}
      813 |      0 |
{"\\x02000000", "\\x433263376364393461613734636530346564373332356139336163
6465623033", "\\x1120323930333435", "\\x133636323632323939"}
      813 |      0 |
{"\\x03000000", "\\x43643365373033636435363532323833336235666261646431343539
6239616130", "\\x1120353438363430", "\\x133239353030333238"}
(5 rows)
```



## vacuum重用的dead tuples空间

```
//
perryn_demo=# vacuum verbose analyze ucheck;
INFO:  vacuuming "public.ucheck"
INFO:  table "ucheck": found 0 removable, 5 nonremovable row versions in 1
out of 1 pages
DETAIL:  0 dead row versions cannot be removed yet, oldest xmin: 4293968109
Skipped 0 pages due to buffer pins, 0 frozen pages.
CPU: user: 0.00 s, system: 0.00 s, elapsed: 0.00 s.
INFO:  vacuuming "pg_toast.pg_toast_16516"
INFO:  table "pg_toast_16516": found 0 removable, 0 nonremovable row
versions in 0 out of 0 pages
DETAIL:  0 dead row versions cannot be removed yet, oldest xmin: 4293968109
Skipped 0 pages due to buffer pins, 0 frozen pages.
CPU: user: 0.00 s, system: 0.00 s, elapsed: 0.00 s.
INFO:  analyzing "public.ucheck"

INFO:  "ucheck": scanned 1 of 1 pages, containing 3 live rows and 0 dead
rows; 3 rows in sample, 3 estimated total rows
VACUUM
```

```
// 会话A插入数据
perryn_demo=# begin;
BEGIN
perryn_demo=# select xmin,xmax,cmin,cmax,* from ucheck;
 xmin | xmax | cmin | cmax | id |                uuid                | code
| md5
-----+-----+-----+-----+----+-----+-----+-----+-----+
--+-----+
    821 |    0 |    5 |    5 |  1 | 6b65bf7e7080ef40110cdae28e145036 |
540085 | 48283533
    821 |    0 |    5 |    5 |  2 | e4f1b77d6b14f9b55f4607b812039074 |
516574 | 67780358
    821 |    0 |    5 |    5 |  3 | 2b81b3a381ec1f70c8f3ddc6af3976b5 |
541362 | 34386270
(3 rows)

perryn_demo=# select xmin,xmax,cmin,cmax,* from ucheck;
 xmin | xmax | cmin | cmax | id |                uuid                | code
| md5
-----+-----+-----+-----+----+-----+-----+-----+-----+
--+-----+
    821 |    0 |    5 |    5 |  1 | 6b65bf7e7080ef40110cdae28e145036 |
540085 | 48283533
    821 | 822 |    0 |    0 |  2 | e4f1b77d6b14f9b55f4607b812039074 |
516574 | 67780358
```

```

      821 |      822 |      0 |      0 |      3 | 2b81b3a381ec1f70c8f3ddc6af3976b5 |
541362 | 34386270
(3 rows)

```

```
perryn_demo=# commit;
```

```
// 会话B 更新数据
```

```
perryn_demo=# begin;
```

```
BEGIN
```

```
perryn_demo=#
```

```
perryn_demo=# select txid_current();
txid_current
```

```
-----
      822
(1 row)
```

```
perryn_demo=# update ucheck set md5=substring(random()::varchar,3,8) where
id>=2;
```

```
UPDATE 2
```

```
perryn_demo=# select xmin,xmax,cmin,cmax,* from ucheck;
```

```

xmin | xmax | cmin | cmax | id |                                uuid                                | code
| md5
-----+-----+-----+-----+-----+-----+-----

```

```

      821 |      0 |      5 |      5 |      1 | 6b65bf7e7080ef40110cdae28e145036 |
540085 | 48283533
      822 |      0 |      0 |      0 |      2 | e4f1b77d6b14f9b55f4607b812039074 |
516574 | 33335208
      822 |      0 |      0 |      0 |      3 | 2b81b3a381ec1f70c8f3ddc6af3976b5 |
541362 | 69940957
(3 rows)

```

```
perryn_demo=# ALTER TABLE ucheck SET (autovacuum_enabled = false,
toast.autovacuum_enabled = false);
```

```
ALTER TABLE
```

```
perryn_demo=# commit;
```

// 未执行vacuum之前的表信息，可以看到t\_xmax=822的有2条记录，这个是会话A插入时候的产生的数据，但是被会话B(事务ID=822)更新数据后,xmax被更新为822.同时会话B插入了2条新的记录，从这里可以看出PG是采用cow策略进行数据的更新

```
perryn_demo=# SELECT t_xmin, t_xmax, tuple_data_split('ucheck'::regclass,
t_data, t_infomask, t_infomask2, t_bits) FROM
heap_page_items(get_raw_page('ucheck', 0));
```

```

t_xmin | t_xmax |
tuple_data_split
-----+-----+-----

```

```

---
      821 |      0 |
{"\x01000000","\x43366236356266376537303830656634303131306364616532386531

```

```

3435303336", "\\x1120353430303835", "\\x133438323833353333"}
      821 |      822 |
{"\\x02000000", "\\x43653466316237376436623134663962353566343630376238313230
3339303734", "\\x1120353136353734", "\\x133637373830333538"}
      821 |      822 |
{"\\x03000000", "\\x43326238316233613338316563316637306338663364646336616633
3937366235", "\\x1120353431333632", "\\x133334333836323730"}
      822 |      0 |
{"\\x02000000", "\\x43653466316237376436623134663962353566343630376238313230
3339303734", "\\x1120353136353734", "\\x133333333335323038"}
      822 |      0 |
{"\\x03000000", "\\x43326238316233613338316563316637306338663364646336616633
3937366235", "\\x1120353431333632", "\\x133639393430393537"}
(5 rows)

```

// 执行vacuum 空间数据被标记清空，但是占用的磁盘并没有归还给操作系统，其从821是会话A的插入事务ID。822是会话B的更新事务的ID,这里有2条空的记录被标记为后面插入数据时候可以被复用。

```

perryn_demo=# vacuum ucheck;
VACUUM
perryn_demo=# SELECT t_xmin, t_xmax, tuple_data_split('ucheck'::regclass,
t_data, t_infomask, t_infomask2, t_bits) FROM
heap_page_items(get_raw_page('ucheck', 0));
 t_xmin | t_xmax |
tuple_data_split
-----+-----+-----
-----
---
      821 |      0 |
{"\\x01000000", "\\x43366236356266376537303830656634303131306364616532386531
3435303336", "\\x1120353430303835", "\\x133438323833353333"}
      |      |
      |      |
      822 |      0 |
{"\\x02000000", "\\x43653466316237376436623134663962353566343630376238313230
3339303734", "\\x1120353136353734", "\\x133333333335323038"}
      822 |      0 |
{"\\x03000000", "\\x43326238316233613338316563316637306338663364646336616633
3937366235", "\\x1120353431333632", "\\x133639393430393537"}
(5 rows)

```

## vacuum full回收的dead tuples空间

```

// 普通vacuum仅仅标记
perryn_demo=# SELECT t_xmin, t_xmax, tuple_data_split('ucheck'::regclass,
t_data, t_infomask, t_infomask2, t_bits) FROM
heap_page_items(get_raw_page('ucheck', 0));
 t_xmin | t_xmax |

```

```

tuple_data_split
-----+-----+-----
-----
---
      821 |      0 |
{"\\x01000000", "\\x43366236356266376537303830656634303131306364616532386531
3435303336", "\\x1120353430303835", "\\x133438323833353333"}
      |      |
      |      |
      822 |      0 |
{"\\x02000000", "\\x43653466316237376436623134663962353566343630376238313230
3339303734", "\\x1120353136353734", "\\x133333333335323038"}
      822 |      0 |
{"\\x03000000", "\\x43326238316233613338316563316637306338663364646336616633
3937366235", "\\x1120353431333632", "\\x133639393430393537"}
(5 rows)

```

// 这里执行vacuum full,可以看出被标记的复用空闲空间归还给操作系统了,但是这个操作会产生表锁。

```

perryn_demo=# vacuum full ucheck;
VACUUM
perryn_demo=# SELECT t_xmin, t_xmax, tuple_data_split('ucheck'::regclass,
t_data, t_infomask, t_infomask2, t_bits) FROM
heap_page_items(get_raw_page('ucheck', 0));
 t_xmin | t_xmax |
tuple_data_split
-----+-----+-----
-----
---
      821 |      0 |
{"\\x01000000", "\\x43366236356266376537303830656634303131306364616532386531
3435303336", "\\x1120353430303835", "\\x133438323833353333"}
      822 |      0 |
{"\\x02000000", "\\x43653466316237376436623134663962353566343630376238313230
3339303734", "\\x1120353136353734", "\\x133333333335323038"}
      822 |      0 |
{"\\x03000000", "\\x43326238316233613338316563316637306338663364646336616633
3937366235", "\\x1120353431333632", "\\x133639393430393537"}
(3 rows)

```